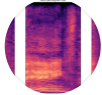


[Open in app](#)

Published in Towards Data Science



mlearnere

[Follow](#)

Sep 17, 2020 · 3 min read · ✨ · 🎧 Listen



Save



Learning from Audio: Wave Forms

An introduction to wave forms and dealing with null data.



Photo by [Jonathan Velasquez](#) on [Unsplash](#)

Related Articles:

- [Learning from Audio: Time Domain Features](#)
- [Learning from Audio: Fou](#)



56



- [Learning from Audio: Spectrograms](#)
- [Learning from Audio: The Mel Scale, Mel Spectrograms, and Mel Frequency Cepstral Coefficients](#)
- [Learning from Audio: Pitch and Chromagrams](#)

Introduction:

Audio is an extremely rich data source. Depending on the `sample rate` — *the number of points sampled per second to quantify the signal* — one second of data could contain thousands of points. Scale this up to hours of recorded audio, and you can see how Machine Learning and Data Science nicely intertwine with signal processing techniques.

This article aims to break down what exactly wave forms are as well as utilize `librosa` in Python for analysis and visualizations — alongside `numpy` and `matplotlib`.

Wave Forms:

`Waves` are repeated `signals` that oscillate and vary in amplitude, depending on their complexity. In the real world, `waves` are continuous and mechanical — which is quite different from computers being discrete and digital.

So, how do we translate something continuous and mechanical into something that is discrete and digital?

This is where the `sample rate` defined earlier comes in. Say, for example, the `sample rate` of the recorded audio is 100. This means that for every recorded second of audio, the computer will place 100 points along the `signal` in attempts to best “trace” the continuous curve. Once all the points are in place, a smooth curve joins them all together for humans to be able to visualize the sound. Since the recorded audio is in terms of `amplitude` and `time`, we can intuitively say that the wave form operates in the `time domain`.

To better understand what something like this sounds like, we will look at three sounds: a kick drum, a guitar, and a snare drum. [The code and data can be found in](#)

my GitHub repository for this article.

```
1 import librosa
2 # Import the wav files. It will be loaded in as a numpy array.
3 # sr represents the samplerate.
4 # Since we did not specify the samplerate when loading in the files,
5 # the default will be set to 22050.
6 guitar, sr = librosa.load('guitar.wav')
7 kick, sr = librosa.load('kick.wav')
8 snare, sr = librosa.load('snare.wav')
```

wavform1.py hosted with ❤ by GitHub

[view raw](#)

Loading in the data.

Now that the data is loaded in, let's visualize these sounds.

```
1 import matplotlib.pyplot as plt
2 #Visualizing waveforms
3 fig, ax = plt.subplots(1,3, figsize = (30,10), sharey = True)
4 librosa.display.waveplot(guitar, sr=sr, ax=ax[0])
5 ax[0].set(title = 'Guitar Waveform')
6 librosa.display.waveplot(kick, sr=sr, ax=ax[1])
7 ax[1].set(title = 'Kick Drum Waveform')
8 librosa.display.waveplot(snare, sr=sr, ax=ax[2])
9 ax[2].set(title = 'Snare Drum Waveform')
10 plt.show()
```

wavform2.py hosted with ❤ by GitHub

[view raw](#)

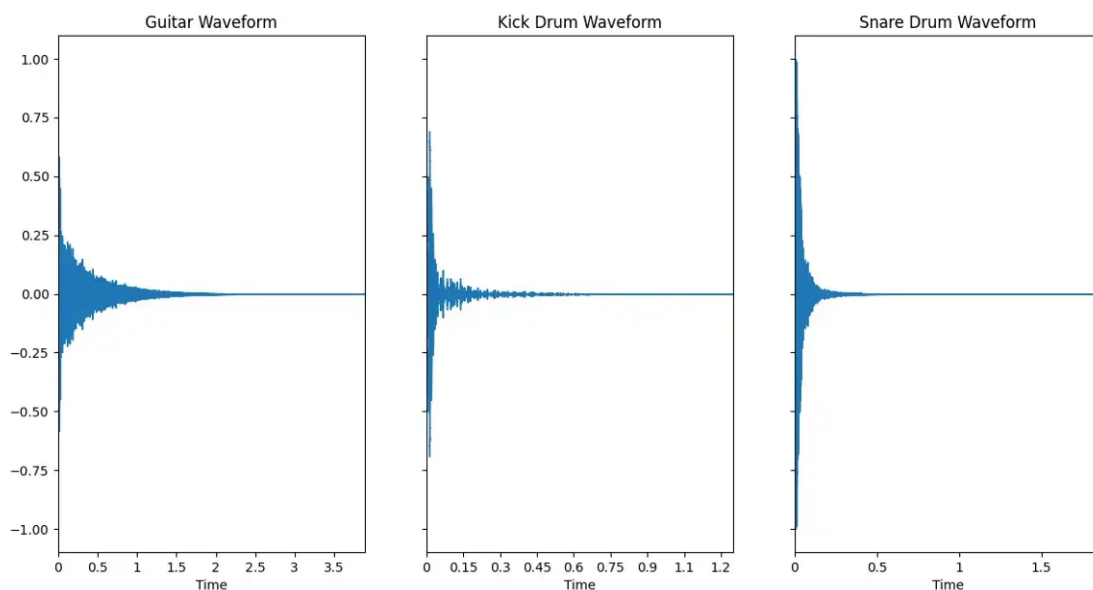


Figure 1

From the get-go, we see some issues with the visualization.

While we can easily tell some differences between the visualizations, it is not as distinct as we would like it to be. We also know that audio signals do not just suddenly disappear, they in fact fade out until it is impossible to perceive. This means that in terms of audio, this constitutes as `null` data.

Null Data in Audio:

There are many ways to treat `null` audio data in the `time domain`. However, this approach often is the simplest.

Given the `signal` and a minimum `threshold` for the `amplitude` of the `signal`:

- Take the `absolute value` of each point in the `signal`
- If the point is greater than the `threshold`, we keep it. Otherwise, we remove it.

```
1 import numpy as np
2 # Create function for mask
3 def env_mask(wav, threshold):
4     # Absolute value
5     wav = np.abs(wav)
6     # Point wise mask determination.
7     mask = wav > threshold
8     return mask
9 # Initialize mask
10 g_mask = env_mask(guitar, sr, 0.005)
11 k_mask = env_mask(kick, sr, 0.005)
12 s_mask = env_mask(snare, sr, 0.005)
13 # Plotting the new signals
14 fig, ax = plt.subplots(1,3, figsize = (30,10), sharey = True)
15 # Visualize wave plots with mask applied.
16 librosa.display.waveplot(guitar[g_mask], sr=sr, ax=ax[0])
17 ax[0].set(title = 'Guitar Waveform')
18 librosa.display.waveplot(kick[k_mask], sr=sr, ax=ax[1])
19 ax[1].set(title = 'Kick Drum Waveform')
20 librosa.display.waveplot(snare[s_mask], sr=sr, ax=ax[2])
21 ax[2].set(title = 'Snare Drum Waveform')
22 plt.show()
```

wavform3.py hosted with ❤ by GitHub

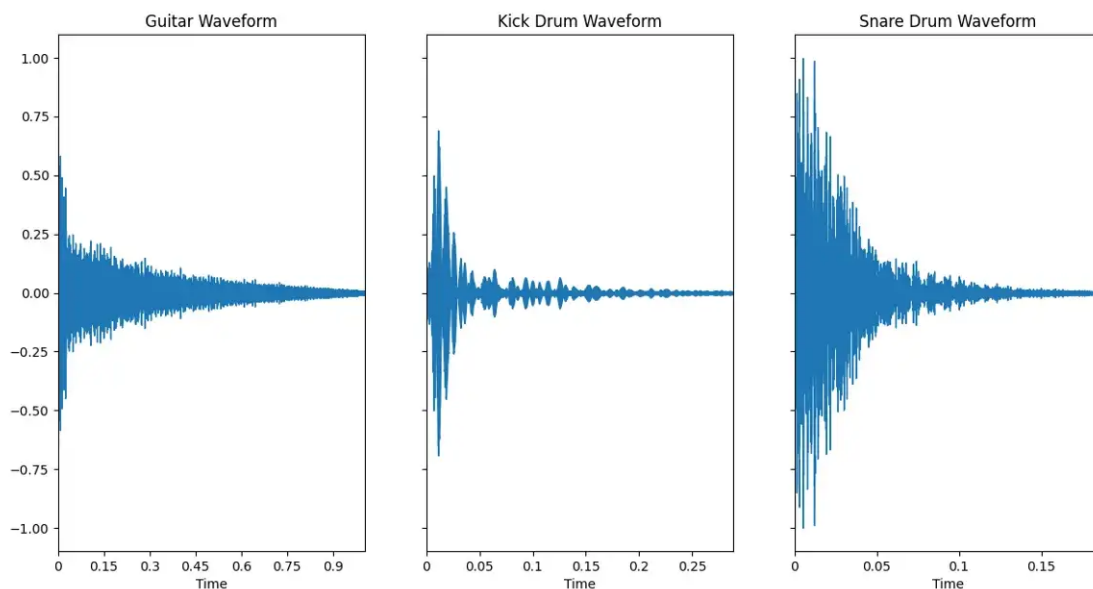
[view raw](#)

Figure 2

You can think of `thresholds` as a sort of parameter for the recordings. Different `thresholds` work differently for various sounds. Playing around with the `threshold` is a good way to see *how* and *why* this visualization changes.

Now that the `null` data has been removed from these recordings, it is much easier to see the personality in each sound. The guitar is much more uniform in shape, drowning out gradually with time. The kick drum hits hard in the beginning and quickly drowns out with some remnants of sound remaining. The snare drum is loud and raucous, something you will not want to listen to repeatedly.

Conclusion:

This concludes the basics of dealing with audio signals in Python with `librosa`. Stay tuned for more articles that delve into more advanced topics of how to learn from audio!

Thank you for reading.

Note: all figures without a source is by the author.

[Audio](#)[Deep Learning](#)[Python](#)[Tutorial](#)[Machine Learning](#)

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

Emails will be sent to shivakmuddam25@gmail.com. [Not you?](#)



Get this newsletter